

```

#Naive Bayes Approach
import pandas as pd
import numpy as np
import nltk
from nltk.corpus import stopwords
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import classification_report, accuracy_score
import os
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import LabelEncoder

train_dataset = 'train.csv'
test_dataset = 'test.csv'

# Check if the path exists
print (os.path.exists(train_dataset))
print (os.path.exists(test_dataset))

True
True

train_df = pd.read_csv(train_dataset, encoding='ISO-8859-1')
test_df = pd.read_csv(test_dataset, encoding='ISO-8859-1')

train_df.head()

{"summary":{"\n  \"name\": \"train_df\", \n  \"rows\": 27481, \n
\n\"fields\": [\n    {\n      \"column\": \"textID\", \n
\n\"properties\": {\n      \"dtype\": \"string\", \n
\n\"num_unique_values\": 27481, \n      \"samples\": [\n
\n\"a7f72a928a\", \n      \"ef42dee96c\", \n      \"07d17131b1\" \n
\n], \n      \"semantic_type\": \"\", \n      \"description\": \"\" \n
\n    }, \n    {\n      \"column\": \"text\", \n      \"properties\":
\n      {\n        \"dtype\": \"string\", \n        \"num_unique_values\":
\n        27480, \n        \"samples\": [\n          \" Enjoy! Family trumps
\n          everything\", \n          \" --of them kinda turns me off of it all.
\n          And then I buy more of them and dig a deeper hole, etc. ;;\", \n
\n          \"Clive it`s my birthday pat me
\n          http://apps.facebook.com/dogbook/profile/view/6386106\" \n        ], \n
\n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
\n      }, \n      {\n        \"column\": \"selected_text\", \n
\n\"properties\": {\n        \"dtype\": \"string\", \n
\n\"num_unique_values\": 22430, \n        \"samples\": [\n
\n\"that is why I drive a (teeny tiny) honda civic\", \n
\n\"Sorry...but, I bet they aren`t that bad...\", \n        \"yummy\" \n
\n], \n        \"semantic_type\": \"\", \n        \"description\": \"\" \n
\n      }, \n      {\n        \"column\": \"sentiment\", \n
\n\"properties\": {\n        \"dtype\": \"category\", \n

```

```

{"num_unique_values": 3, "samples": [{"neutral", "negative", "positive"}, {"semantic_type": "", "description": ""}], {"column": "Time of Tweet", "properties": {"dtype": "category", "num_unique_values": 3, "samples": [{"morning", "noon", "night"}], {"semantic_type": "", "description": ""}], {"column": "Age of User", "properties": {"dtype": "category", "num_unique_values": 6, "samples": [{"0-20", "21-30", "70-100"}], {"semantic_type": "", "description": ""}], {"column": "Country", "properties": {"dtype": "category", "num_unique_values": 195, "samples": [{"Belgium", "Sierra Leone"}], {"semantic_type": "", "description": ""}], {"column": "Population -2020", "properties": {"dtype": "number", "std": 150494590, "min": 801, "max": 1439323776, "num_unique_values": 195, "samples": [{"109581078, 11589623, 7976983}], {"semantic_type": "", "description": ""}], {"column": "Land Area (Km\u00b2)", "properties": {"dtype": "number", "std": 1807424.6900064405, "min": 0.0, "max": 16376870.0, "num_unique_values": 193, "samples": [{"2267050.0, 1280000.0, 100250.0}], {"semantic_type": "", "description": ""}], {"column": "Density (P/Km\u00b2)", "properties": {"dtype": "number", "std": 2013, "min": 2, "max": 26337, "num_unique_values": 136, "samples": [{"71, 331, 400}], {"semantic_type": "", "description": ""}]}], {"type": "dataframe", "variable_name": "train_df"}

```

```
train_df.tail()
```

```

{"summary": {"name": "train_df", "rows": 5, "fields": [{"column": "textID", "properties": {"dtype": "string", "num_unique_values": 5, "samples": [{"4f4c4fc327", "6f7127d9d7", "f67aae2310"}], {"semantic_type": "", "description": ""}], {"column": "text", "properties": {"dtype": "string", "num_unique_values": 5, "samples": [{"I've wondered about rake to. The client has made it clear .NET only, don't force devs to learn a new lang"}]}]}]}

```

```

#agile #ccnet\", \n          \"    All this flirting going on - The ATG
smiles. Yay. ((hugs))\", \n          \" Yay good for both of you.
Enjoy the break - you probably need it after such hectic weekend Take
care hun xxxx\", \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          }, \n          { \n          \"column\":
\"selected_text\", \n          \"properties\": { \n          \"dtype\":
\"string\", \n          \"num_unique_values\": 5, \n          \"samples\":
[ \n          \", don't force\", \n          \"All this flirting going
on - The ATG smiles. Yay. ((hugs))\", \n          \"Yay good for both
of you.\" \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          }, \n          { \n          \"column\":
\"sentiment\", \n          \"properties\": { \n          \"dtype\":
\"string\", \n          \"num_unique_values\": 3, \n          \"samples\":
[ \n          \"negative\", \n          \"positive\", \n
\"neutral\" \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          }, \n          { \n          \"column\":
\"Time of Tweet\", \n          \"properties\": { \n          \"dtype\":
\"string\", \n          \"num_unique_values\": 3, \n          \"samples\":
[ \n          \"night\", \n          \"morning\", \n          \"noon\" \n
], \n          \"semantic_type\": \"\", \n          \"description\": \"\", \n
} \n          }, \n          { \n          \"column\": \"Age of User\", \n
\"properties\": { \n          \"dtype\": \"string\", \n
\"num_unique_values\": 5, \n          \"samples\": [ \n          \"46-
60\", \n          \"0-20\", \n          \"60-70\" \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\", \n
} \n          }, \n          { \n          \"column\": \"Country\", \n
\"properties\": { \n          \"dtype\": \"string\", \n
\"num_unique_values\": 5, \n          \"samples\": [ \n          \"Greece\", \n
\"Guinea\", \n
\"Grenada\" \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          }, \n          { \n          \"column\":
\"Population -2020\", \n          \"properties\": { \n          \"dtype\":
\"number\", \n          \"std\": 11311650, \n          \"min\": 112523, \n
\"max\": 31072940, \n          \"num_unique_values\": 5, \n
\"samples\": [ \n          10423054, \n          13132795, \n
112523 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          }, \n          { \n          \"column\":
\"Land Area (Km\u00b2)\", \n          \"properties\": { \n
\"dtype\": \"number\", \n          \"std\": 99481.25712916981, \n
\"min\": 340.0, \n          \"max\": 246000.0, \n
\"num_unique_values\": 5, \n          \"samples\": [ \n
128900.0, \n          246000.0, \n          340.0 \n          ], \n
\"semantic_type\": \"\", \n          \"description\": \"\", \n
} \n          }, \n          { \n          \"column\": \"Density (P/Km\u00b2)\", \n
\"properties\": { \n          \"dtype\": \"number\", \n          \"std\":
108, \n          \"min\": 53, \n          \"max\": 331, \n
\"num_unique_values\": 5, \n          \"samples\": [ \n          81, \n
53, \n          331 \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\", \n          } \n          } \n          ] \n          }, \"type\": \"dataframe\"}

```

```
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27481 entries, 0 to 27480
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   textID                 27481 non-null  object
1   text                   27480 non-null  object
2   selected_text          27480 non-null  object
3   sentiment               27481 non-null  object
4   Time of Tweet          27481 non-null  object
5   Age of User             27481 non-null  object
6   Country                 27481 non-null  object
7   Population -2020        27481 non-null  int64
8   Land Area (Km²)         27481 non-null  float64
9   Density (P/Km²)         27481 non-null  int64
dtypes: float64(1), int64(2), object(7)
memory usage: 2.1+ MB
```

```
test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4815 entries, 0 to 4814
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   textID                 3534 non-null  object
1   text                   3534 non-null  object
2   sentiment               3534 non-null  object
3   Time of Tweet          3534 non-null  object
4   Age of User             3534 non-null  object
5   Country                 3534 non-null  object
6   Population -2020        3534 non-null  float64
7   Land Area (Km²)         3534 non-null  float64
8   Density (P/Km²)         3534 non-null  float64
dtypes: float64(3), object(6)
memory usage: 338.7+ KB
```

```
train_df.isnull().sum()
```

```
textID      0
text        1
selected_text  1
sentiment    0
Time of Tweet  0
Age of User  0
Country      0
Population -2020  0
Land Area (Km²)  0
Density (P/Km²)  0
dtype: int64
```

```

train_df = train_df.dropna()
train_df.isnull().sum()

textID      0
text         0
selected_text 0
sentiment    0
Time of Tweet 0
Age of User  0
Country      0
Population -2020 0
Land Area (Km²) 0
Density (P/Km²) 0
dtype: int64

test_df.isnull().sum()

textID      1281
text         1281
sentiment    1281
Time of Tweet 1281
Age of User  1281
Country      1281
Population -2020 1281
Land Area (Km²) 1281
Density (P/Km²) 1281
dtype: int64

test_df = test_df.dropna()
test_df.isnull().sum()

textID      0
text         0
sentiment    0
Time of Tweet 0
Age of User  0
Country      0
Population -2020 0
Land Area (Km²) 0
Density (P/Km²) 0
dtype: int64

nltk.download('stopwords')

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Unzipping corpora/stopwords.zip.

True

def preprocess_text(text):
    # Convert to lowercase

```

```

text = text.lower()
# Remove stopwords
stop_words = set(stopwords.words('english'))
text = " ".join([word for word in text.split() if word not in
stop_words])
return text

train_df['processed_text'] = train_df['text'].apply(preprocess_text)
train_df.head()

{"summary":{"\n  \"name\": \"train_df\",\n  \"rows\": 27480,\n  \"fields\": [\n    {\n      \"column\": \"textID\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 27480,\n        \"samples\": [\n          \"6c5505a37c\",\n          \"126b1e6a22\",\n          \"5bc4e623c4\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"text\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 27480,\n        \"samples\": [\n          \" Enjoy! Family trumps everything\",\n          \" --of them kinda turns me off of it all. And then I buy more of them and dig a deeper hole, etc. ;;\",\n          \"Clive it`s my birthday pat me http://apps.facebook.com/dogbook/profile/view/6386106\",\n          \"\",\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        ]\n      },\n      {\n        \"column\": \"selected_text\",\n        \"properties\": {\n          \"dtype\": \"string\",\n          \"num_unique_values\": 22430,\n          \"samples\": [\n            \"that is why I drive a (teeny tiny) honda civic\",\n            \"Sorry...but, I bet they aren`t that bad...\",\n            \"yummy\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"sentiment\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 3,\n          \"samples\": [\n            \"neutral\",\n            \"negative\",\n            \"positive\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"Time of Tweet\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 3,\n          \"samples\": [\n            \"morning\",\n            \"noon\",\n            \"night\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"Age of User\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 6,\n          \"samples\": [\n            \"0-20\",\n            \"21-30\",\n            \"70-100\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      },\n      {\n        \"column\": \"Country\",\n        \"properties\": {\n          \"dtype\": \"category\",\n          \"num_unique_values\": 195,\n          \"samples\": [\n            \"Philippines\",\n            \"Belgium\",\n            \"Sierra Leone\"\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        }\n      }\n    ]\n  }\n}

```

```

n      },\n      {\n          \"column\": \"Population -2020\", \n          \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 150497157, \n              \"min\": 801, \n              \"max\": 1439323776, \n              \"num_unique_values\": 195, \n              \"samples\": [\n                  109581078, \n                  11589623, \n                  7976983\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n          }\n      },\n      {\n          \"column\": \"Land Area (Km\\u00b2)\", \n          \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 1807457.3166921895, \n              \"min\": 0.0, \n              \"max\": 16376870.0, \n              \"num_unique_values\": 193, \n              \"samples\": [\n                  2267050.0, \n                  1280000.0, \n                  100250.0\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n          }\n      },\n      {\n          \"column\": \"Density (P/Km\\u00b2)\", \n          \"properties\": {\n              \"dtype\": \"number\", \n              \"std\": 2013, \n              \"min\": 2, \n              \"max\": 26337, \n              \"num_unique_values\": 136, \n              \"samples\": [\n                  71, \n                  331, \n                  400\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n          }\n      },\n      {\n          \"column\": \"processed_text\", \n          \"properties\": {\n              \"dtype\": \"string\", \n              \"num_unique_values\": 27268, \n              \"samples\": [\n                  \"thought like really hot. room hot sleep\", \n                  \"praying get better soon sweet one , sorry still well\", \n                  \"damm feel like song dead gone travis garland\"\n              ], \n              \"semantic_type\": \"\", \n              \"description\": \"\"\n          }\n      }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"train_df\"}

```

```

test_df['processed_text'] = test_df['text'].apply(preprocess_text)
test_df.head()

```

```

{"summary": "{\n  \"name\": \"test_df\", \n  \"rows\": 3534, \n  \"fields\": [\n    {\n      \"column\": \"textID\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3534, \n        \"samples\": [\n          \"142108215\", \n          \"fb08563a7b\", \n          \"9a2c6ae21c\"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"text\", \n      \"properties\": {\n        \"dtype\": \"string\", \n        \"num_unique_values\": 3534, \n        \"samples\": [\n          \"Thank you so much phaoloo !!!!\", \n          \"Midnight ice-cream weather! So **** bored\", \n          \"Ohh i forgot to tell you last night that when i was a alton towers i touched a shark it was amazing !!!! it was nt a massive one tho\"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"sentiment\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"neutral\", \n          \"positive\", \n          \"negative\"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Time of Tweet\", \n      \"properties\": {\n        \"dtype\": \"category\", \n        \"num_unique_values\": 3, \n        \"samples\": [\n          \"\", \n          \"\", \n          \"\"\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\": \"dataframe\", \"variable_name\": \"test_df\"}

```


```
[
    \n          \n"morning",\n          \n"noon",\n          \n"night"\n
],\n    \n    \n"semantic_type": \n",\n    \n    \n"description": \n"\n
}\n    },\n    {\n        \n"column": \n"Age of User",\n        \n"properties": {\n            \n"dtype": \n"category",\n            \n"num_unique_values": 6,\n            \n"samples": [\n                \n"0-20",\n                \n"21-30",\n                \n"70-100"\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    },\n    {\n        \n"column": \n"Country",\n        \n"properties": {\n            \n"dtype": \n"category",\n            \n"num_unique_values": 195,\n            \n"samples": [\n                \n"Philippines",\n                \n"Belgium",\n                \n"Sierra Leone"\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    },\n    {\n        \n"column": \n"Population -2020",\n        \n"properties": {\n            \n"dtype": \n"number",\n            \n"std": 146875664.36244348,\n            \n"min": 801.0,\n            \n"max": 1439323776.0,\n            \n"num_unique_values": 195,\n            \n"samples": [\n                109581078.0,\n                11589623.0,\n                7976983.0\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    },\n    {\n        \n"column": \n"Land Area (Km\u00b2)",\n        \n"properties": {\n            \n"dtype": \n"number",\n            \n"std": 1839133.911427382,\n            \n"min": 0.0,\n            \n"max": 16376870.0,\n            \n"num_unique_values": 193,\n            \n"samples": [\n                2267050.0,\n                1280000.0,\n                100250.0\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    },\n    {\n        \n"column": \n"Density (P/Km\u00b2)",\n        \n"properties": {\n            \n"dtype": \n"number",\n            \n"std": 1967.012367010644,\n            \n"min": 2.0,\n            \n"max": 26337.0,\n            \n"num_unique_values": 136,\n            \n"samples": [\n                400.0,\n                71.0,\n                331.0\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    },\n    {\n        \n"column": \n"processed_text",\n        \n"properties": {\n            \n"dtype": \n"string",\n            \n"num_unique_values": 3527,\n            \n"samples": [\n                \n"thank much phao! !!!",\n                \n"i'm glad you're little prissy well. it's obvious much love w/the treatment she's getting",\n                \n"3d version sold regular version is!"\n            ],\n            \n"semantic_type": \n",\n            \n"description": \n"\n        }\n    }\n]\n}", "type": "dataframe", "variable_name": "test_df"}
```

```
sns.countplot(x='sentiment', data=train_df, palette=['red', 'green', 'blue'])
plt.title("Distribution of Sentiment Labels")
plt.show()
```

<ipython-input-17-a42c8626a2cf>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.


```
sns.countplot(x='sentiment', data=train_df, palette=['red', 'green', 'blue'])
```



```
vectorizer = TfidfVectorizer(max_features=5000)
train_TFIDF = vectorizer.fit_transform(train_df['processed_text'])
test_TFIDF = vectorizer.transform(test_df['processed_text'])

# Assign X and y for the training data
X = train_TFIDF
y = train_df['sentiment']

NB_model = MultinomialNB()
NB_model.fit(X, y)

MultinomialNB()

predictions = NB_model.predict(test_TFIDF)

print(f"Accuracy: {accuracy_score(test_df['sentiment'],
predictions)}")
print(classification_report(test_df['sentiment'], predictions))
```

Accuracy: 0.6386530843237125

	precision	recall	f1-score	support
negative	0.73	0.50	0.59	1001
neutral	0.55	0.77	0.65	1430
positive	0.76	0.60	0.67	1103
accuracy			0.64	3534
macro avg	0.68	0.62	0.63	3534
weighted avg	0.67	0.64	0.64	3534

```
input_text = ["What a bad product!"]
```

```
input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")
```

Predicted sentiment: negative

```
input_text = ["Last session of the day"]
```

```
input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")
```

Predicted sentiment: neutral

```
input_text = ["I hate meetings!"]
```

```
input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")
```

Predicted sentiment: negative

```
input_text = ["I love cars!"]
```

```
input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")
```

Predicted sentiment: positive

```
input_text = ["Do you want help?"]
```

```
input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")
```

Predicted sentiment: neutral

```

input_text = ["You are good girl"]


input_TFIDF = vectorizer.transform(input_text)
predicted_sentiment = NB_model.predict(input_TFIDF)
print(f"Predicted sentiment: {predicted_sentiment[0]}")

Predicted sentiment: positive

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
cm = confusion_matrix(test_df['sentiment'], predictions)

disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=NB_model.classes_)
disp.plot()
plt.show()

```



```

import joblib

# Assuming you have already trained your model and vectorizer
joblib.dump(NB_model, 'NB_model.pkl')      # Save the Naive Bayes
model
joblib.dump(vectorizer, 'vectorizer.pkl')   # Save the TF-IDF
vectorizer

```

```
['vectorizer.pkl']  
from google.colab import files  
  
files.download('NB_model.pkl')  
files.download('vectorizer.pkl')  
<IPython.core.display.Javascript object>  
<IPython.core.display.Javascript object>  
<IPython.core.display.Javascript object>  
<IPython.core.display.Javascript object>
```